

RB Weekly AI Brief - Issue 23 - 08.09.2026

Covering the week of 08.09.2026 · Issue 23 of the RB Weekly AI Brief

Recurring themes: Regulatory & HTA Signals (3 of last 4 issues) · Regulation & Policy (3 of last 4 issues) · Healthcare & Life Sciences (3 of last 4 issues) · Models & Research (3 of last 4 issues)

AI News Roundup

Regulatory & HTA Signals

NICE Publishes AI Delivery Plan for NHS 2026/27

[IN CASE YOU MISSED IT] On 6 August 2026, NICE published a new delivery plan setting out how it intends to support the safe, effective, and scalable adoption of artificial intelligence across the NHS. The plan focuses on three areas: using AI for evidence generation and evaluation, assessing AI-powered health technologies, and deploying AI to improve NICE's own internal processes. By March 2027, NICE intends to publish an AI best-practice methods framework covering the use of AI in evidence generation, synthesis, and review — and has set a target of halving staff time required to produce each piece of guidance by 2030.

***So what?** Read this as a statement of intent, not a new evidentiary bar: NICE's own binding methods framework isn't due until March 2027, and an ISPOR 2026 poster found AI methods were reported in just three of 1,391 real HTA submission documents reviewed across NICE, CDA-AMC, and IQWiG — all three from NICE itself, under its existing 2024 position statement. Pharma and HEOR teams should track the March 2027 date, but shouldn't mistake this plan for evidence that AI submissions are already being scrutinised differently in practice.*

NICE

Regulation & Policy

MHRA National Commission Misses Self-Imposed Summer Deadline for AI Healthcare Recommendations

[ONGOING] The UK's MHRA-hosted National Commission into the Regulation of AI in Healthcare has been working since October 2025 to develop a new regulatory framework for AI products in healthcare. Its evidence-gathering phase is complete: in June 2026, the MHRA published two landmark reports drawing on input from 760 individuals and organisations, patients, clinicians, and industry. But the Commission's own recommendations — originally due 'by summer 2026' per its founding terms of reference — remain unpublished as of early September, with the MHRA's CEO confirming the agency will act on them once they land.

***So what?** Life sciences and health-tech companies preparing regulatory and market access strategies for AI-enabled products in the UK are now waiting on a single, overdue deliverable rather than an evolving process — the evidence base has been public since June, so the gap is specifically in the Commission's own decision-making, worth monitoring closely as it will directly shape MHRA's updated framework for clinical AI and software as a medical device.*

MHRA / GOV.UK

Healthcare & Life Sciences

FDA TEMPO Pilot Lets Generative-AI Medical Devices Reach Patients Pre-Authorization

The FDA launched the TEMPO (Technology-Enabled Meaningful Patient Outcomes) pilot, provisionally allowing generative AI medical devices to reach the market before formal marketing authorization while regulators gather real-world performance data. Four devices have been accepted, including products from Cadence and Limbic, with the pilot tied directly to the Medicare ACCESS model — a CMS Innovation Center experiment in paying for chronic disease management technology. The program gives FDA a live testing ground for regulating care-influencing AI tools rather than waiting for a fully settled regulatory framework.

***So what?** TEMPO signals that real-world evidence collected under enforcement discretion — not pre-market RCT data alone — may shape the evidentiary bar for future AI medical device coverage decisions under Medicare, directly informing how HEOR and HTA teams design evidence generation strategies for AI-enabled therapeutics and digital health submissions.*

FDA

Nobel Laureate David Baker Launches \$95M AI BioDesign Lab for Protein & Gene Design

The Allen Institute, University of Washington, and Fred Hutchinson Cancer Center jointly launched AI BioDesign on September 3 — a five-year, \$95 million initiative funded by the Paul Allen estate that pairs bench scientists with machine-learning teams to generate massive biological datasets and train AI models to design novel proteins and genes from scratch. Led by 2024 Nobel Chemistry laureate David Baker, the project aims to produce new disease therapies, plastic-dissolving proteins, and programmed cells. Early experimental campaigns have already processed millions of genetic sequences with dedicated ML teams directing subsequent assay designs.

***So what?** AI-designed proteins and genes moving from computational hypothesis to experimental validation at this scale and speed will compress the target-identification and lead-generation timelines that underpin pharma R&D investment cases, requiring HEOR and market access teams to adapt economic models to assets with shorter, AI-accelerated discovery histories and novel mechanism-of-action evidence packages.*

Allen Institute

Models & Research

OpenEvidence Launches Medical AI Model Family Including Darwin, First to Score 100% on MedQA

OpenEvidence released a new family of medical AI models on September 3, 2026, making three — Osler (~5s), Sackett (~30s), and Snow (~5 min) — available free to all verified clinicians, with a fourth model, Darwin, held in research-only preview. Darwin is described by OpenEvidence as the first AI in history to achieve a perfect 100% on MedQA, the leading independent benchmark of medical AI, and also leads on MedXpertQA (72.8%), HealthBench Professional (82.7%), and NOHARM (87.2%). **Caveat:** per independent reporting, Darwin's perfect score was achieved on a reduced, physician-re-annotated set of 660 questions, not the original 1,273-question MedQA test split — a self-curated benchmark, not the standard one other models have been scored against. Darwin is restricted to institutional partners and academic researchers due to dual-use concerns in virology, immunology, and human genetics.

***So what?** A clinician-facing AI platform used by over 1.1 million NPI-verified U.S. clinicians that claims benchmark-leading accuracy on medical evidence synthesis could accelerate the real-world adoption of evidence-based prescribing patterns — but the self-modified benchmark underlying its headline claim is a reminder for pharma and market access teams to treat vendor-reported "firsts" as marketing until independently reproduced on the standard test.*

OpenEvidence

OpenAI Chief Scientist Calls for Voluntary AI Research Slowdown, Citing Alignment Gaps

On September 6, OpenAI Chief Scientist Jakub Pachocki published an essay titled 'An Alien Mind' arguing that no AI lab has solved alignment and monitoring to a sufficient degree to continue scaling at maximum speed, and that current chain-of-thought monitoring capabilities are progressively diminishing in reliability. Pachocki called for voluntary slowdowns to become 'commonplace' until shared safety standards are established, and urged governments to prioritize international coordination on AI development. The essay was published just three days after OpenAI shipped GPT-6 Astra — its most powerful model to date — creating a striking internal contradiction at the world's leading AI lab.

***So what?** If voluntary AI slowdowns or mandatory safety audits gain traction across frontier labs, pharma and HTA organizations that have built AI-accelerated drug development, clinical evidence synthesis, or regulatory submission timelines into their near-term pipeline plans will need to stress-test those timelines against a landscape where frontier model capability growth may be deliberately paced.*

OpenAI

Academic Paper Summaries

Selected from PubMed · Published within the last 12 months · New selections each week

Domain Paper — HEOR / Health Economics / Market Access

Value of AI in Critical Care Using Real-World Evidence on Intensive Care Unit Mortality Prediction: Cost-Utility Analysis.

Lee GM, Won JY, Cho EY, Kim JH, Chung KS, Kim KJ, Yun SJ, Lee HJ, Kim JH. · Journal of medical Internet research · 2026

#HTA · #MarketAccess · #RealWorldEvidence

This study built a lifetime Markov model to evaluate the cost-utility of an AI-based mortality prediction tool activated during ICU admission in Korea, using national claims data and published utility estimates. The AI-assisted strategy produced an ICER of approximately \$6,200 per QALY — well below Korea's societal willingness-to-pay threshold of roughly \$29,600 per QALY — with an 84% probability of being cost-effective at that threshold in probabilistic sensitivity analysis. For HEOR and market access teams, this is a genuine formal cost-utility analysis (ICERs, QALYs, PSA) applied specifically to an AI clinical tool, rather than a diagnostic-accuracy study dressed up as economic evidence.

PMID: 42612087

[PubMed →](#)

[DOI →](#)

AI Research Paper 1

ChatGPT in urogynecology: Comparing large language model responses to human experts.

Rotem R, Simon C, Rottenstreich M, O'reilly BA, Weintraub AY, O'Sullivan OE. · Acta obstetrica et gynecologica Scandinavica · 2025

#ClinicalAI · #PatientOutcomes · #Diagnostics

This study asked 203 urogynecology patients to rate responses to common clinical questions answered by either ChatGPT or a specialist consultant, finding that patients scored the AI higher on clarity, helpfulness, and reassurance across most questions. The results suggest AI chat tools could serve as a valuable first point of contact for patients who feel stigma around discussing pelvic floor conditions. However, the study used a single consultant comparator and exploratory design, so findings should be interpreted cautiously before informing deployment decisions.

PMID: 41128396

[PubMed →](#)

[DOI →](#)

AI Research Paper 2

Comparative performance of large language models for patient-initiated ophthalmology consultations.

Huang M, Wang X, Zhou S, Cui X, Zhang Z, Xu Y, Yang W, Chi W. · Frontiers in public health · 2025

#ClinicalAI · #PatientOutcomes · #Diagnostics

Researchers tested five leading large language models against 31 common patient ophthalmology questions, rating each on accuracy, safety, and accessibility. ChatGPT-4o and DeepSeek-V3 outperformed the others overall, while different models showed strengths suited to different user types, from lay patients to those with medical training. The findings point toward a potential role for AI in pre-consultation triage, though randomized trials are still needed to confirm real-world patient benefit.

PMID: 41059182

[PubMed →](#)

[DOI →](#)

This brief is produced using an AI-assisted workflow with expert human editorial review. Stories are curated and sources verified before publication. Readers are encouraged to consult source links directly.