AI LITERACY SERIES

What Happens to Your Data When You Use an LLM?

A plain-English guide to file uploads, chat content, and why it matters more than you think

Ryan Bishop | March 2026 | ryanbishop.co.uk

The question nobody asks — until it matters

Most people using AI tools have never stopped to ask a simple question: **where does my data actually go?** The assumption is that a conversation with an LLM is private — a bit like thinking out loud. It isn't.

This piece covers two distinct but related scenarios: uploading a file to an LLM, and typing sensitive information directly into a chat window. The risks are real, they vary significantly by platform, and most organisations have no policy in place to manage either.

To make this concrete, we'll use pharmaceutical pricing and clinical data as our worked example — a high-stakes context where the consequences of data exposure are not abstract.

Scenario 1: Uploading a file

Imagine a market access professional uploads an unpublished clinical trial summary to ChatGPT to help draft a payer dossier. The file contains pre-publication endpoints and proprietary study design.

When a file is uploaded, it is loaded into the model's context window — a temporary working memory for that session. The model processes it, responds, and the session ends. But "temporary" is relative. What happens next depends entirely on the platform, the account tier, and the settings in place.

The practical risk is not technical leakage. Your pre-publication data will not appear verbatim in someone else's output. LLMs do not work that way architecturally. The real risks are more mundane and, in some ways, more serious:

- **Data retained on provider servers** for days, months, or years depending on platform and tier
- **Potential use in model training** if you are on a consumer account and have not opted out
- **Server-side access** by provider staff for safety review, legal compliance, or in the event of a data breach
- **Regulatory and contractual breach** — uploading trial data almost certainly violates ICH E6(R3) GCP obligations, sponsor data confidentiality requirements, and likely your employment contract

The jurisdiction risk is separate — and significant. If you use DeepSeek, that file is now on servers in mainland China, subject to the Chinese National Intelligence Law. That is not a theoretical concern; it is explicitly stated in their privacy policy.

Scenario 2: Typing it into the chat window

Now imagine a pricing analyst types the following into a chat window: "PRODUCT X is listed at £X. Our internal target is £Y. Help me structure a negotiation argument."

This is arguably the higher-risk behaviour — not because the architecture is different, but because people type things into chat windows that they would never put in an email. It feels informal. Ephemeral. It isn't.

Chat content follows exactly the same data pathway as file uploads: processed server-side, subject to retention policies, potentially eligible for training data. There are two additional dimensions worth understanding:

The informal behaviour problem. Employees routinely input commercially sensitive pricing strategy, product names, internal targets, and negotiation positions into consumer-tier AI tools without any awareness of where that information goes. The chat window is not a private notepad.

The aggregate inference risk. If enough users input similar data — drug names, price points, negotiation positions — the model may develop calibrated responses that reflect aggregate knowledge. This is difficult to prove and speculative in nature, but it is architecturally plausible and worth naming.

Platform by platform: what the policies actually say

All five major platforms have been verified against their current published privacy policies. The picture varies — and the consumer versus enterprise distinction is critical across every provider.

Platform	Tier	Training use?	Retention	Key note
ChatGPT	Consumer	Yes, opt-out available	30 days*	Default is ON. Enterprise excluded by default.
Claude	Consumer (Free/Pro/Max)	Yes from Sep 2025, opt-out available	30 days (5 yrs if opted in)	Policy shifted Aug 2025. Work/Enterprise excluded.
Copilot	Consumer (personal)	Yes, opt-out available	18 months	Uploaded files not used for training. Enterprise excluded.
Gemini	Free/unpaid	Yes — incl. human review	Up to 3 years	Google's own hub states: 'don't enter confidential info'.
DeepSeek	All tiers	Yes — NO opt-out	Indefinite	Data stored in China. Subject to Chinese national law. Multiple EU bans.

*ChatGPT note: Some deleted chat logs being retained longer under ongoing US litigation (as of early 2026). Enterprise and API users excluded from training by default across all platforms listed above.

The governance gap

The question most organisations need to answer is not technical. It is behavioural and organisational:

- Do employees know what constitutes sensitive data in this context?
- Do they know which tools are approved — and which are not?
- Is there a clear data classification policy that extends to AI tool usage?
- Is there any training that addresses the informal chat window behaviour — not just file uploads?

Does your organisation have clear answers to these questions? If not, that is the gap to close — and in pharmaceutical and healthcare contexts specifically, where data carries regulatory weight and commercial sensitivity, it is a gap with real consequences.

⚠️ Key Takeaways & What To Do

For individuals

- **Check your settings now.** On every platform you use, locate the training opt-out toggle and turn it off.
- **Never use consumer-tier tools for sensitive data.** Free ChatGPT, personal Copilot, free Gemini, and DeepSeek are not appropriate for unpublished data, pricing information, or patient-level data.
- **Avoid DeepSeek for professional use entirely.** No opt-out, indefinite retention, servers in mainland China, multiple regulatory bans across EU countries.

For organisations

- **Build a data classification policy that includes AI tools.** Define what constitutes sensitive data and which tools are approved for which data types.
- **Deploy enterprise-tier tools, not consumer accounts.** Microsoft 365 Copilot, Claude for Work, ChatGPT Enterprise, and Google Workspace Gemini all exclude data from training by default.
- **Train people — not just on policy, but on the 'why'.** Understanding that the chat window is not private changes behaviour far more effectively than a policy document.

The bottom line

AI tools are not inherently unsafe. But using consumer-tier tools with sensitive data, without understanding what happens to it, is a governance failure — not a technology failure. The information to make better decisions exists. Using it is a choice.

Want to discuss what responsible AI implementation looks like for your organisation? **Get in touch: ryan@ryanbishop.co.uk**

Sources & verification

All platform statements verified against published privacy policies (February–March 2026):

- OpenAI: openai.com/policies — training opt-out via Settings › Data Controls
- Anthropic: anthropic.com/news/updates-to-our-consumer-terms — policy update August 2025
- Microsoft: learn.microsoft.com/copilot/microsoft-365/microsoft-365-copilot-privacy
- Google Gemini: support.google.com/gemini/answer/13594961 — Gemini Apps Privacy Hub
- DeepSeek: cdn.deepseek.com — Privacy Policy (updated 22 December 2025); China National Intelligence Law (2017), Article 7
- EU regulatory actions: IAPP, Lexology, DPO Centre — DeepSeek investigations and bans (2025)